

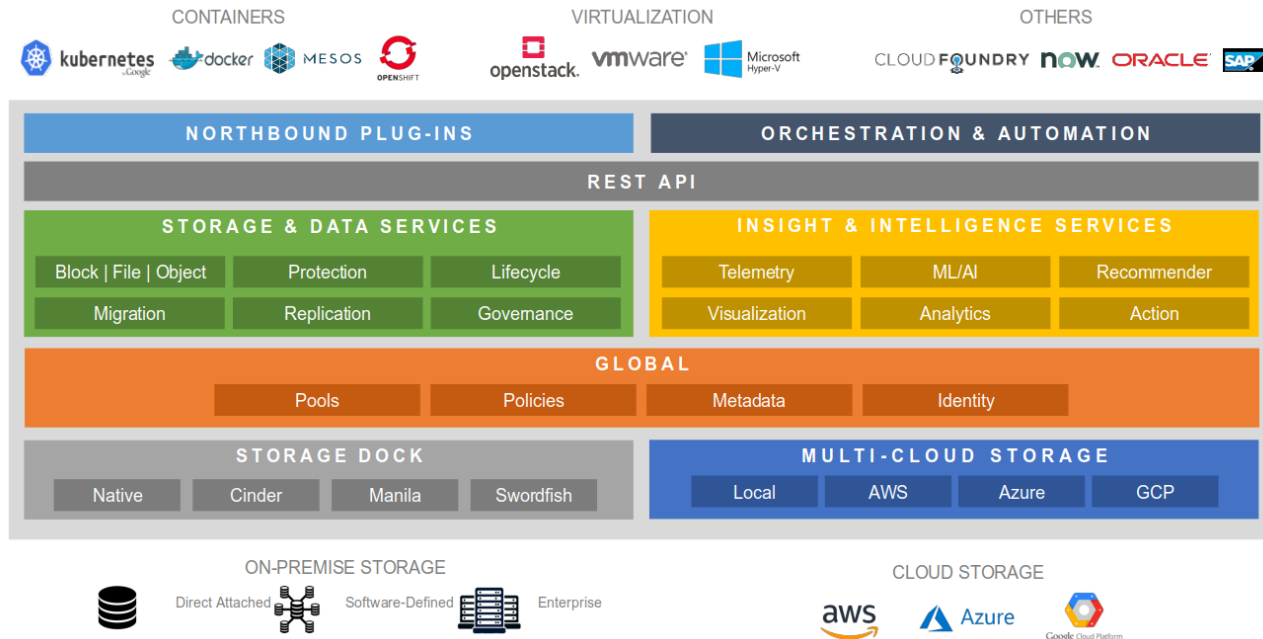


# The OpenSDS POC UseCase



# OpenSDS Architecture and Project

## OpenSDS: Open Intelligent Data Life Cycle Management



### ❖ SUSHI

#### The Northbound Plug-ins Project

Common plug-ins to enable OpenSDS storage services for cloud and application frameworks

### ❖ HOTPOT

#### The Storage Controller Project

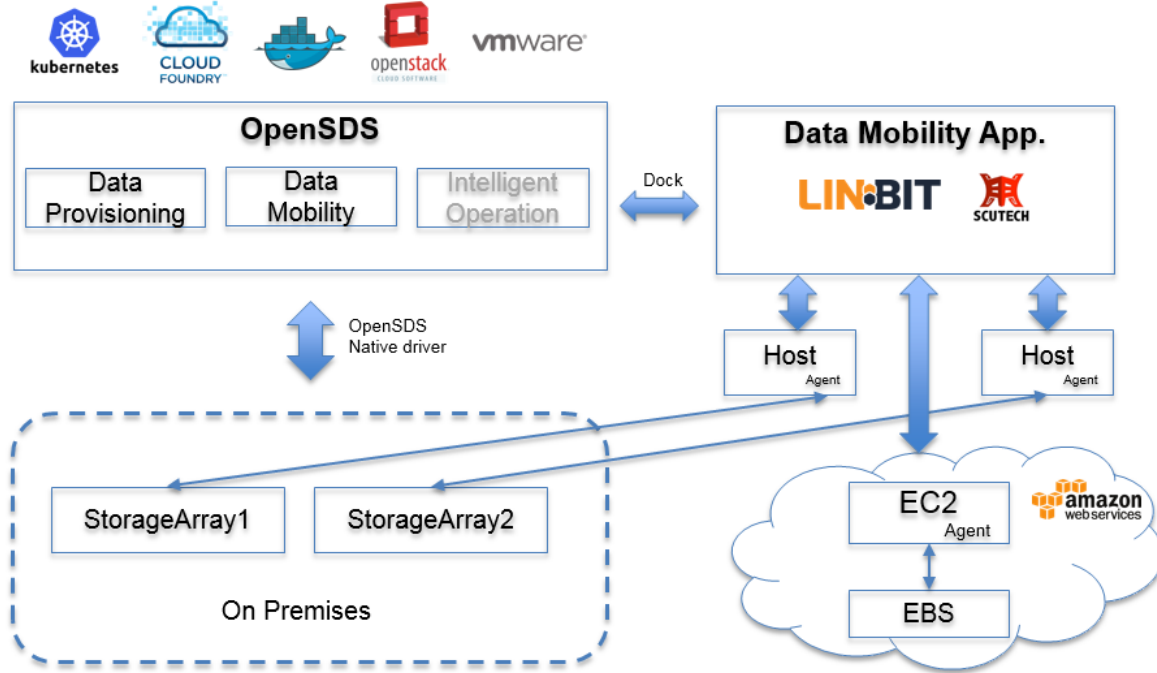
Single control for block, file, and object services across storage on premise and in clouds

### ❖ Gelato

#### The Multi-Cloud Project

Policy based multi-cloud data control to enable data mobility across clouds

# OpenSDS Use Cases



## Data Provisioning

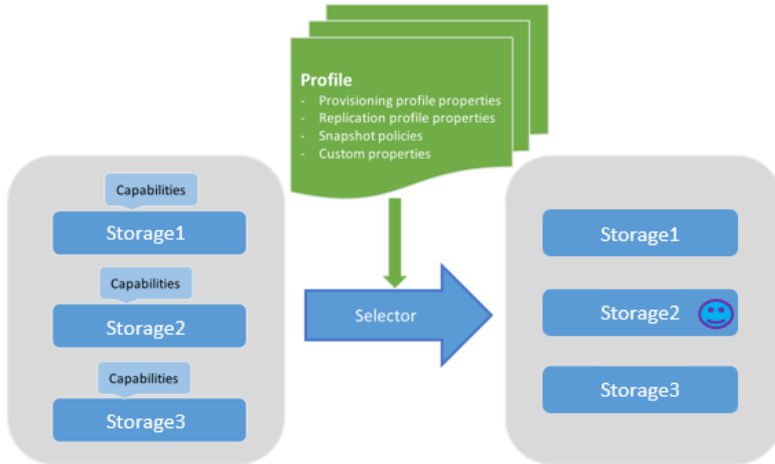
1. Intelligent Storage Profile
2. Container Storage Service
3. Multi-Apps. Data Service

## Data Mobility

1. Data Protection in Cloud
2. Array-based Replication
3. Host-based Replication
4. Multi-Cloud Data Mobility

# 1. Intelligent Storage Profile

## Profile Capabilities Mapping



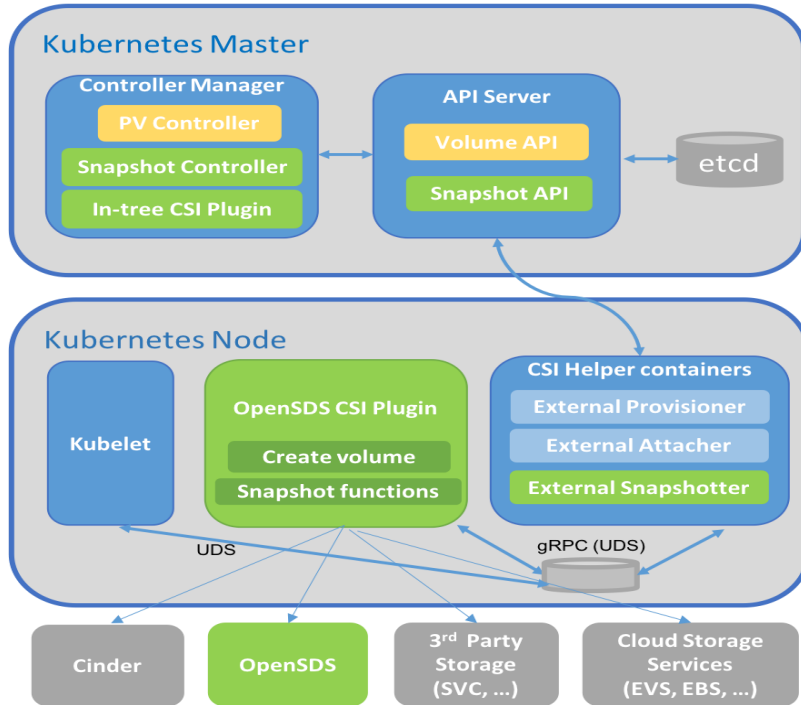
### Use case steps :

1. Discovery and display storage arrays:
  - Storage1 : SSD, IOPS:800465, FC
  - Storage2 : NL-SAS, IOPS:7000, iSCSI
2. Create profile with some capabilities
  - Name: webapp-profile
  - AccessProtocol: iSCSI
  - DiskType: NL-SAS
3. Apply for volume based on “webapp-profile”:
  - Name: webapp-volume
  - Size: 10G
  - Profile: webapp-profile

### Expected results :

Select Storage2 to create a volume named “webapp-volume”.

# Container Storage Service ( CSI )



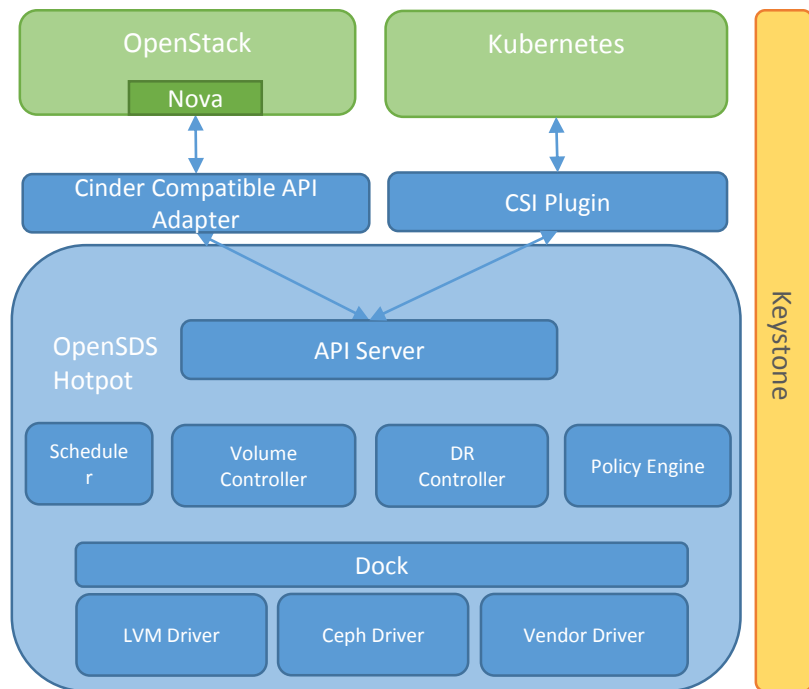
## Use case steps :

1. Create profile with some parameters:
  - Name: webapp-profile
  - AccessProtocol: iSCSI
  - DiskType: NL-SAS
2. Create Storage Class in kubernetes:
  - StorageClass:
    - Name: webapp-sc
    - Provisioner: csi-opensdsplugin
    - Parameters:
      - Profile: webapp-profile
3. Create a pod named “nginx” whose PVC is based on the storage class “webapp-sc”.

## Expected results :

A volume is allocated from OpenSDS and attached to “nginx” pod.

# Multi-Apps. Data Service(OpenStack & Container)



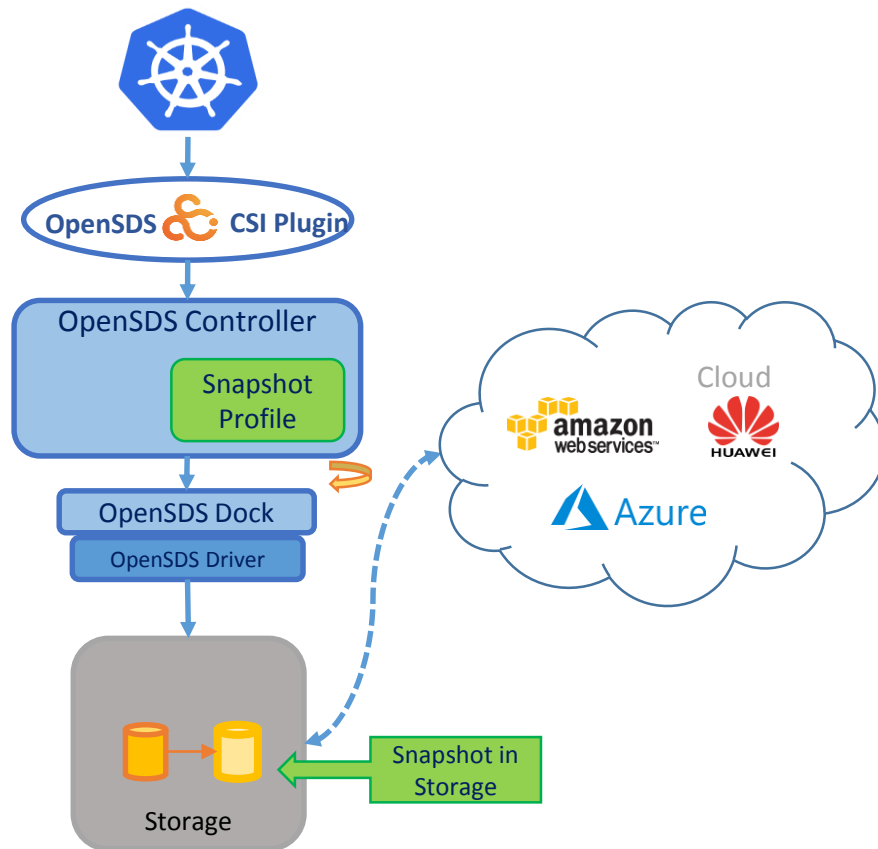
## Use case steps :

1. Provisioning for OpenStack:
  - a. Create volume "my-openstack-volume" using Cinder Command.
  - b. Create Nova instance and attach volume "my-openstack-volume" to it.
2. Provisioning for Kubernetes:
  - a. Create Storage Class in kubernetes:
    - StorageClass:
      - Name: webapp-sc
      - Provisioner: csi-opensdsplugin
      - Parameters:
        - Profile: webapp-profile
  - b. Create a pod named "nginx" whose PVC is based on the storage class "webapp-sc".

## Expected results :

1. A volume is allocated from OpenSDS and attached to OpenStack nova instance.
2. A volume is allocated from OpenSDS and attached to Kubernetes pod.

# Data Protection in Multi-Cloud



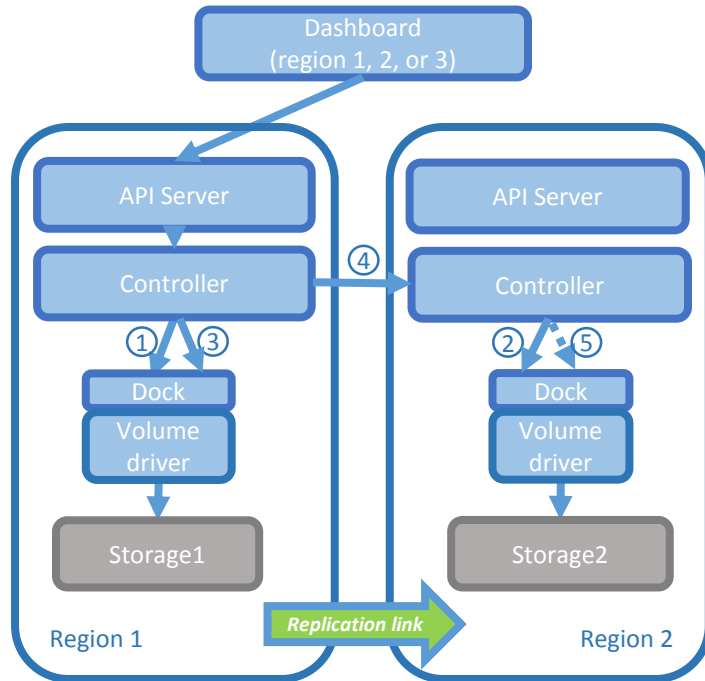
## Use case steps :

1. Create Storage Class in kubernetes:
  - StorageClass:
    - Name: webapp-sc
    - Provisioner: csi-opensdsplugin
    - Parameters:
      - Profile: webapp-profile
2. Create Snapshot Class in Kubernetes
  - VolumeSnapshotClass:
    - Name: csi-opensds-snapclass
    - Snapshotter: csi-opensdsplugin
    - Parameters:
      - Profile: upload-snapshot-profile
3. Create a pod named "nginx" whose PVC is based on the storage class "webapp-sc".
4. Create snapshot for PVC:
  - a. OpenSDS will create a snapshot
  - b. OpenSDS uploads snapshot to cloud.
5. Create PVC from snapshot:
  - a. OpenSDS downloads snapshot from cloud
  - b. OpenSDS create a volume from snapshot

## Expected results :

1. A volume is allocated from OpenSDS and attached to pod.
2. A volume is recovered from snapshot across cloud.

# Array-Based replication Scenario



## Scenario: Homogeneous Storage Replication

### Use case steps:

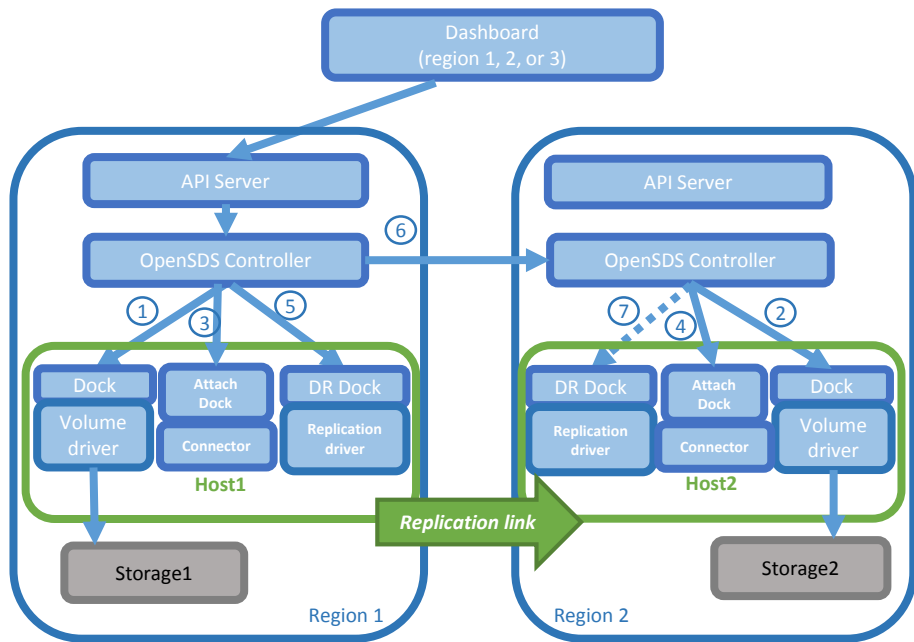
1. Creates Volume
  - a. Creates source volume
  - b. Controller 2 create target volume
2. Create Replication
  - a. Creates source replication
  - b. Controller 2 creates entry in db
3. Write "Hello World" in source volume
4. Execute replication

### Expected result:

1. There is "Hello World" in target volume



# Host-Based replication Scenario



## Scenario: heterogeneous Storage Replication

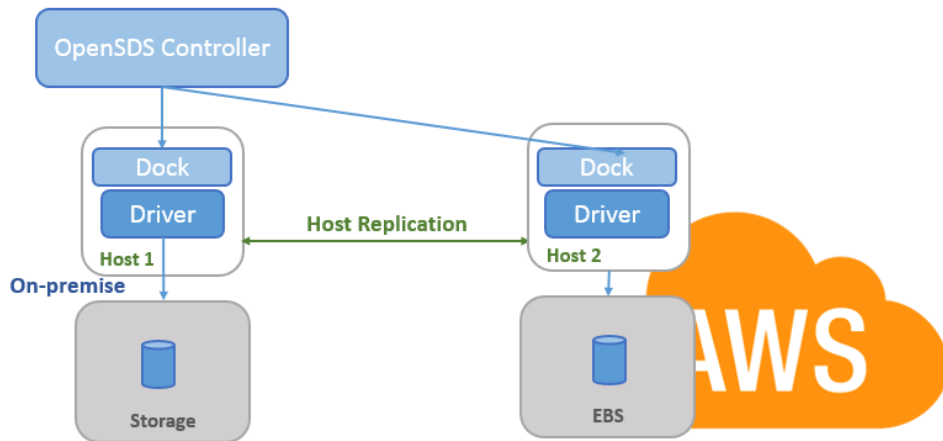
### Use case steps:

1. Create Volume
  - a. Controller 1 creates source volume
  - b. Controller 2 creates target volume
2. Create Replication
  - a. Attach source volume to Host 1
  - b. Attach target volume to Host 2
  - c. Controller 1 creates source replication
  - d. Controller 2 creates entry in db
3. Write "Hello World" in source volume

### Expected result:

1. There is "Hello World" in secondary volume

# Multi-Cloud Data Mobility



**Scenario: Data flow from On-premise to public cloud(AWS)**

## Use case steps:

1. Create Volume
  - a. Create source volume in OpenSDS
  - b. Attach source volume to Host 1
  - c. Create target EBS volume
  - d. Create an ECS and attach target volume to it
2. Create Replication across cloud
  - a. Controller 1 creates source replication
  - b. Controller 2 creates entry in db
3. Write "Hello World" in source volume

## Expected result:

1. There is "Hello World" in target volume



# Thank You

@opensds\_io

#opensds

