

# Managing and Protecting Persistent Volumes for Kubernetes

Xing Yang, Principal Architect, Huawei



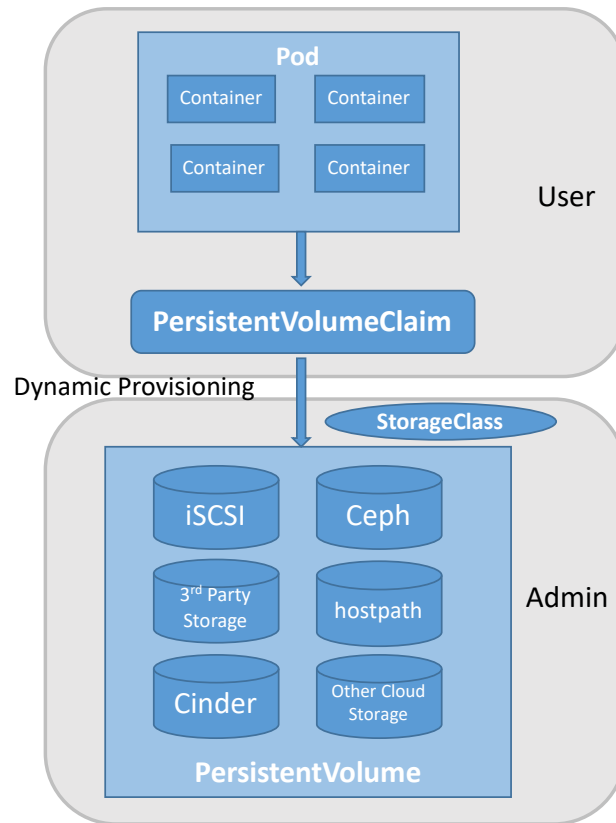
# Agenda

- Kubernetes Persistent Volumes
- Container Storage Interface (CSI)
- Provision and Manage Persistent Volumes using OpenSDS
  - OpenSDS Architecture
  - Mapping OpenSDS Profile to K8S StorageClass
  - Policy Driven SPDM
- Data Protection for Persistent Volumes
- Disaster Recovery for Persistent Volumes
  - Array-based Replication
  - Host-based Replication
- Demo

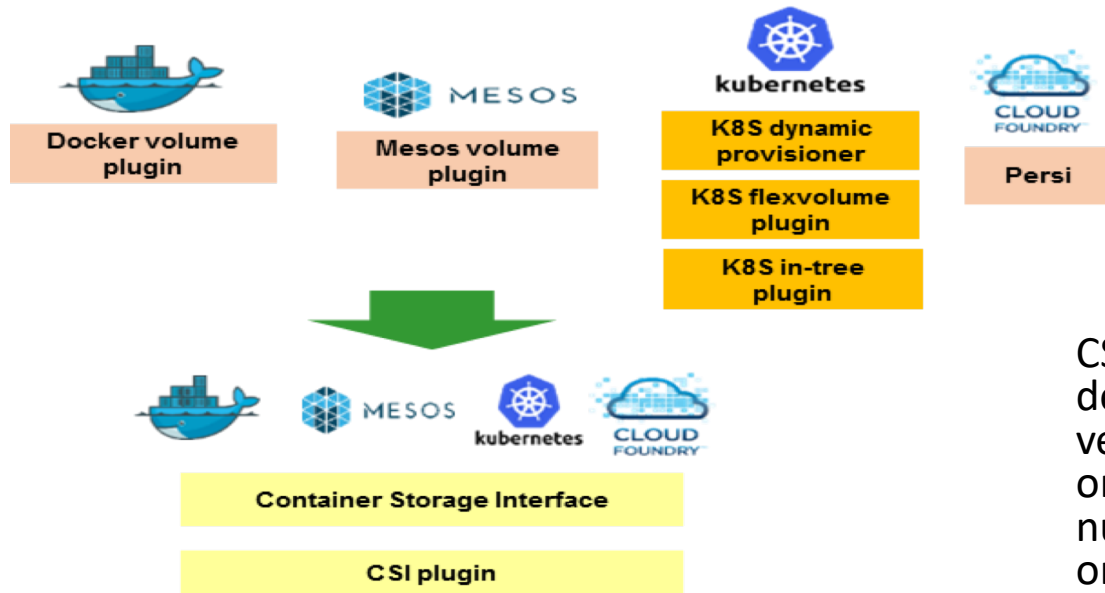


# Kubernetes Persistent Volumes

- A PersistentVolume (PV) is a piece of storage in the cluster that has been provisioned by an administrator.
- A PersistentVolumeClaim (PVC) is a request for storage by a user through a StorageClass.
- A StorageClass provides a way for administrators to describe the “classes” of storage they offer. Different classes might map to different quality-of-service levels (or “profiles”) in other storage systems.
- A StorageClass needs to specify a provisioner for dynamic provisioning.

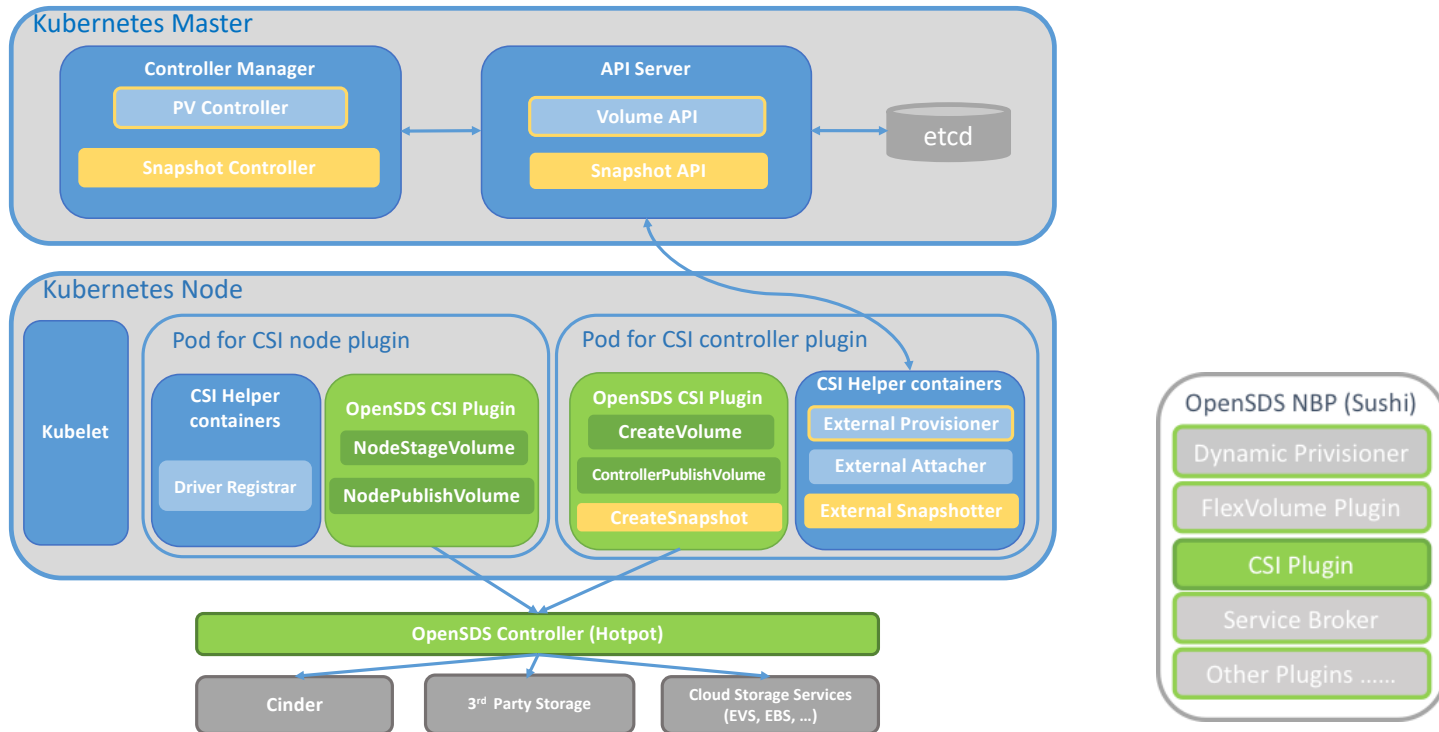


# Container Storage Interface (CSI)

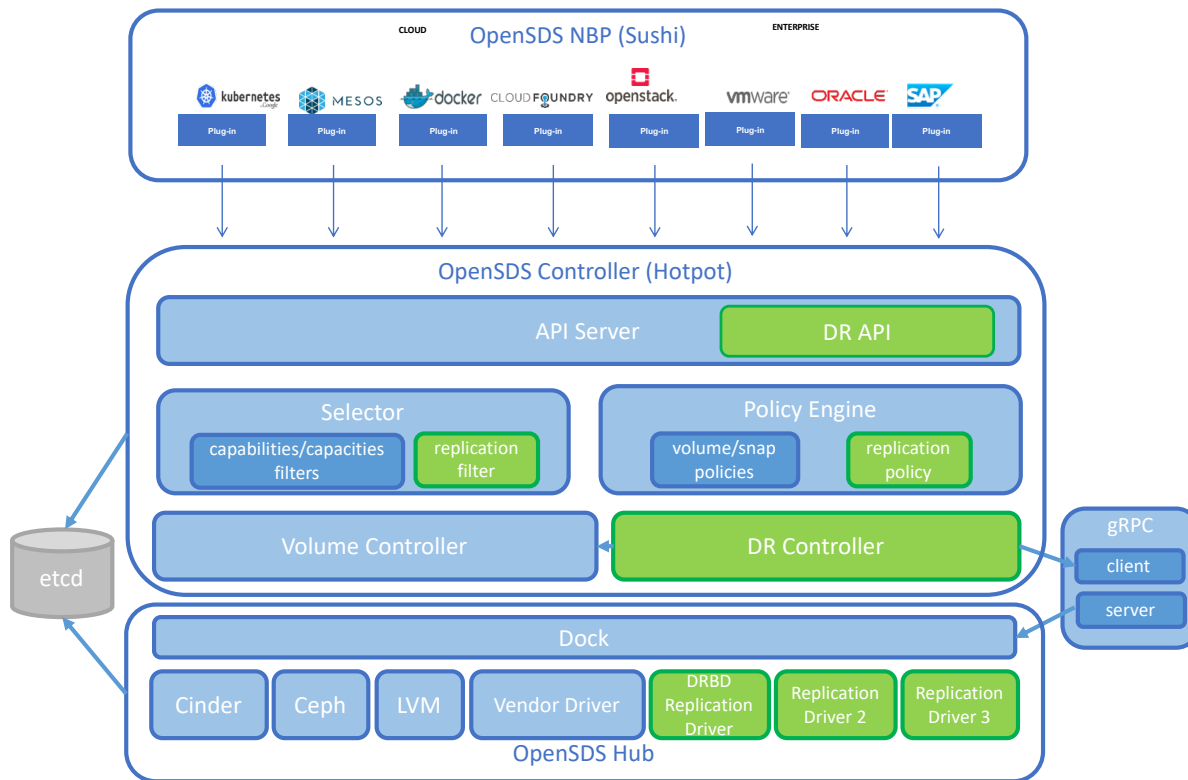


CSI is an industry standard defined to enable storage vendors to develop a plugin once and have it work across a number of container orchestration systems.

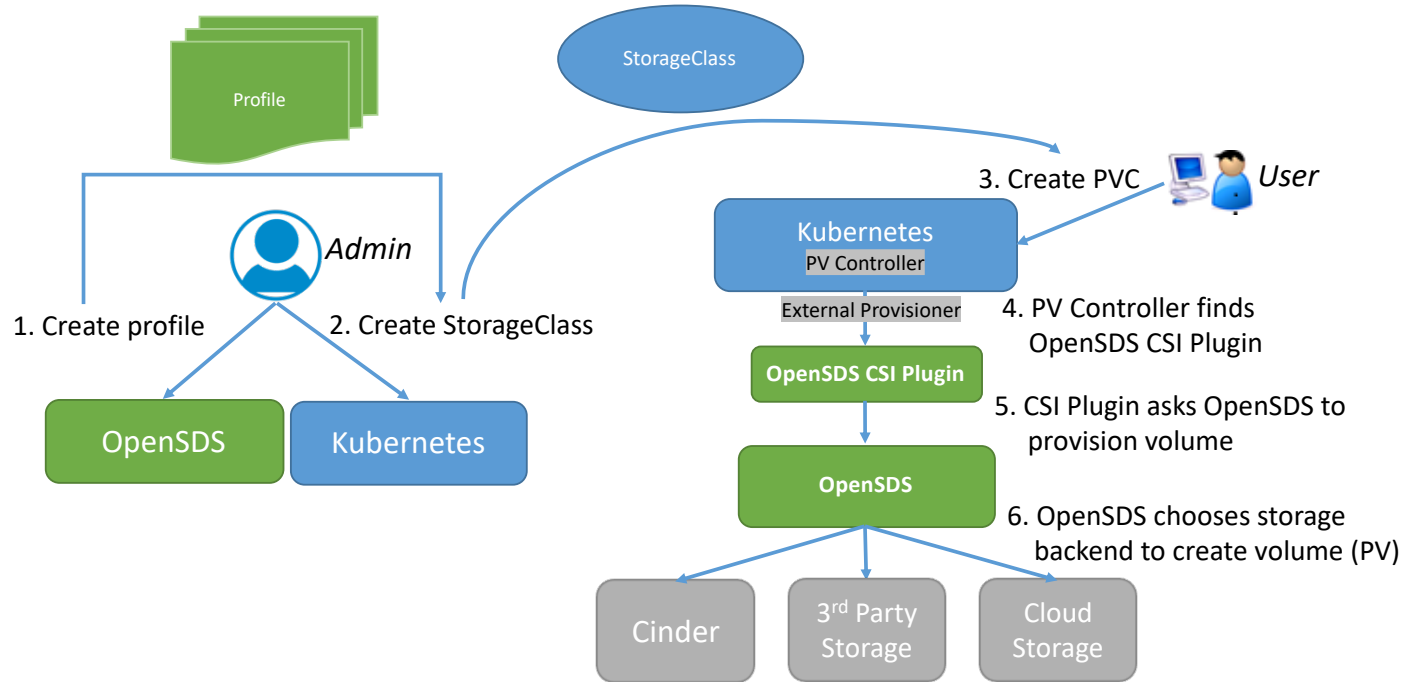
# Provision and Manage Persistent Volumes using OpenSDS



# OpenSDS Architecture



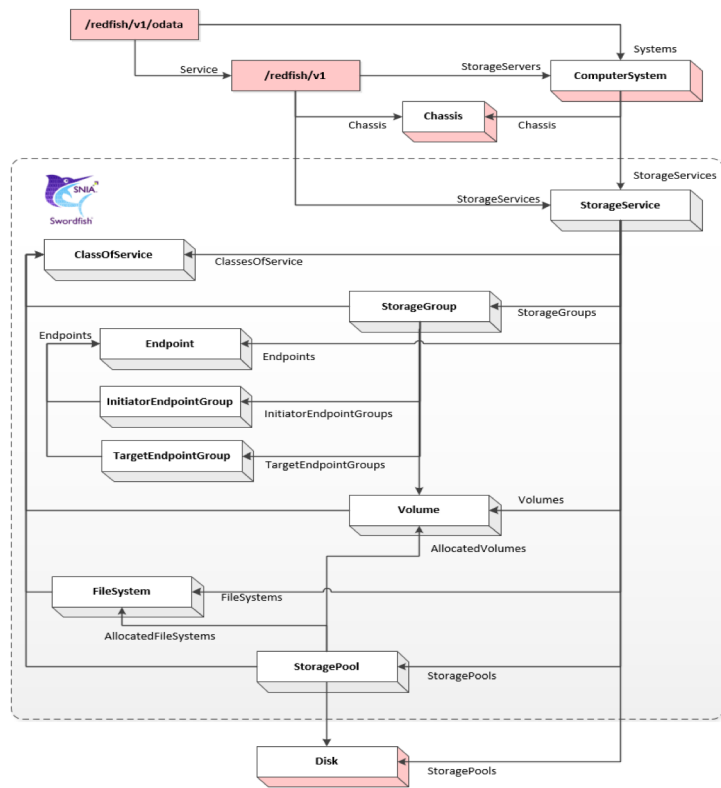
# Mapping OpenSDS Profile to K8S StorageClass



# Policy Driven SPDM



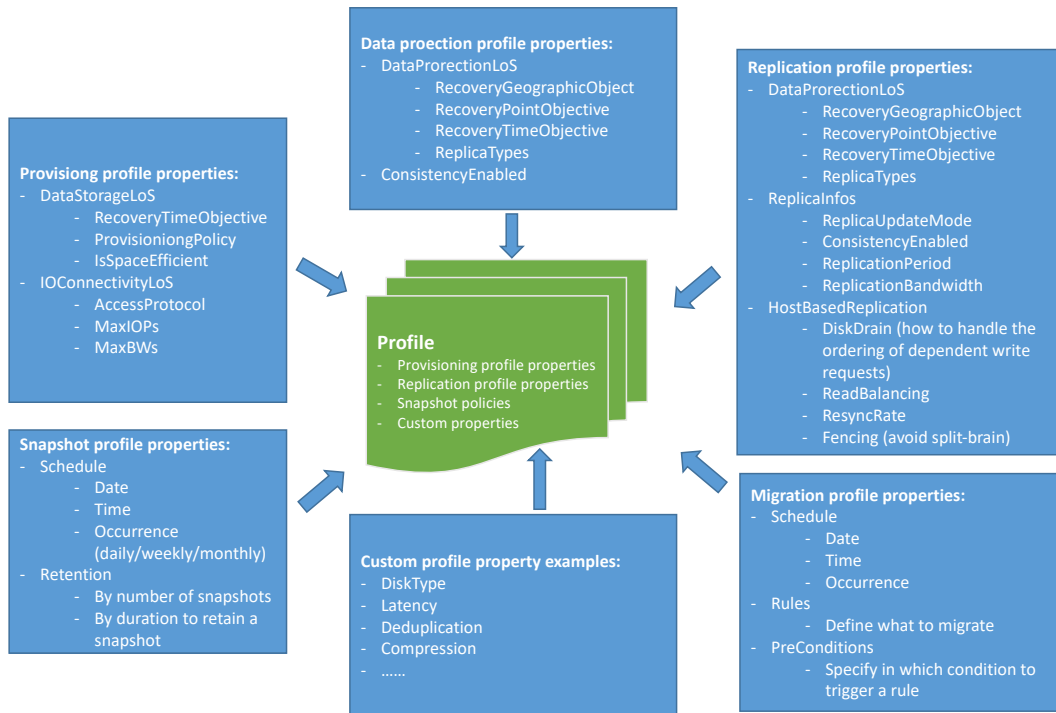
- OpenSDS profile is based on Swordfish specification.
- The SNIA Swordfish™ specification helps to provide a unified approach for the management of storage and servers in hyperscale and cloud infrastructure environments, supported by multiple storage vendors.
- An extension of the DMTF (Distributed Management Task Force) Redfish specification.
  - Redfish is designed by the DMTF's Scalable Platforms Management Forum (SPMF) to create and publish an open industry standard specification and schema for management of scalable platform hardware. It is a RESTful interface over HTTPS in JSON format based on OData v4.



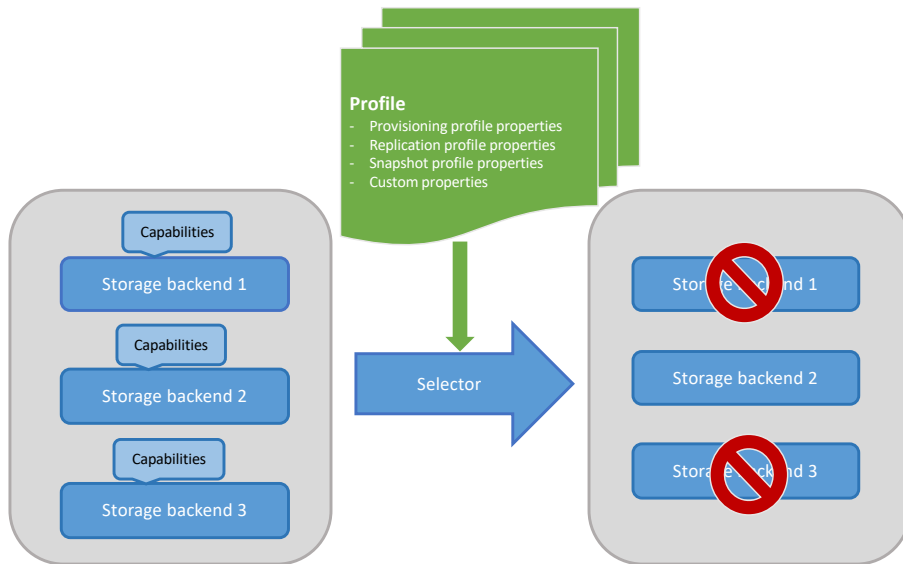
Source: Swordfish\_v1.0.5\_Specification



# Profile Definitions



# Mapping Profiles to Capabilities



# Profile Example

## DATA STORAGE

- **DataStorageLoS**
  - RecoveryTimeObjective (Immediate, Nearline, Offline, Online)
  - ProvisioningPolicy (thin, thick)
  - IsSpaceEfficient (true, false)

## DATA PERFORMANCE



- **IOConnectivitLoS**
  - AccessProtocol (iSCSI, FC, RBD ...)
  - MaxIOPS
  - MaxBWS

## REPLICATION

- **DataProtectionLoS**
  - RecoveryGeographicObjective
  - RecoveryTimeObjective
  - RecoveryPointObjective
  - ReplicaTypes
- **ReplicaInfos**
  - ReplicaitonUpdateMode
  - ConsistencyEnabled
  - ReplicationPeriod
  - ReplicationBandwidth

## SNAPSHOT

- **Schedule**
  - Date
  - Time
  - Occurrence (daily/weekly/monthly)
- **Retention**
  - By number of snapshots
  - By duration to retain a snapshot



# StorageClass with Profile Parameter

## HighPerformanceSC.yaml

```

apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: opensds-csi-high-performance-sc
provisioner: csi-opensdsplugin
parameters:
  profile: High-Performance
  
```

## HighPerformancePVC.yaml

```

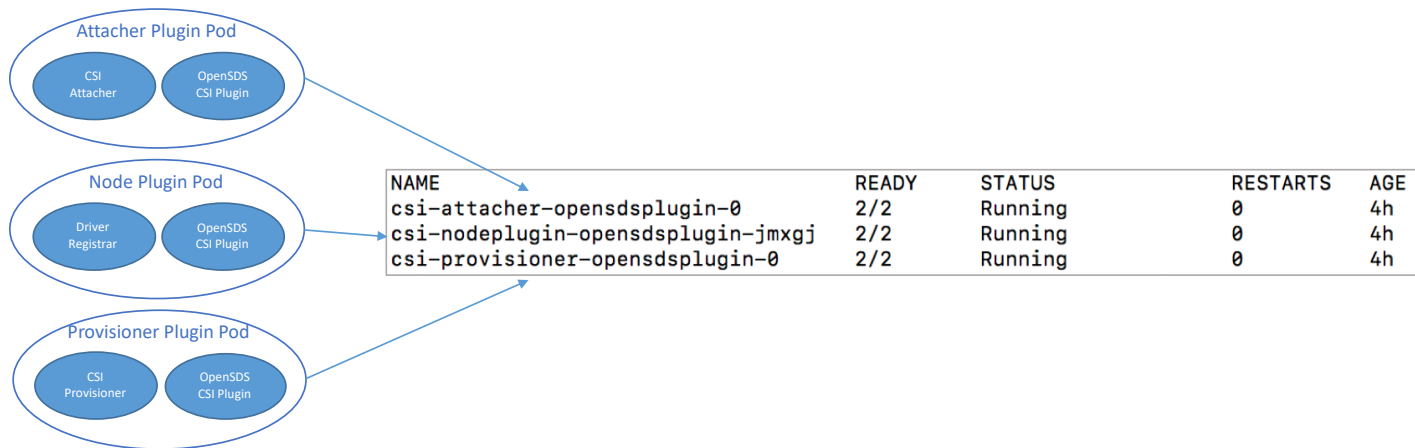
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: opensds-csi-high-performance-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
  storageClassName: opensds-csi-high-performance-sc
  
```

*Note: profile parameter can be profile id or name*



# Running OpenSDS CSI Plugin

- Create OpenSDS CSI plugin pods:  
`kubectl create -f csi/server/deploy/kubernetes`
- Three pods can be found by `kubectl get pod`:



# Using OpenSDS Volume

- Create nginx application

```
kubectl create -f
```

```
csi/server/examples/kubernetes/nginx.yaml
```

- An OpenSDS volume is mounted at */var/lib/www/html*.

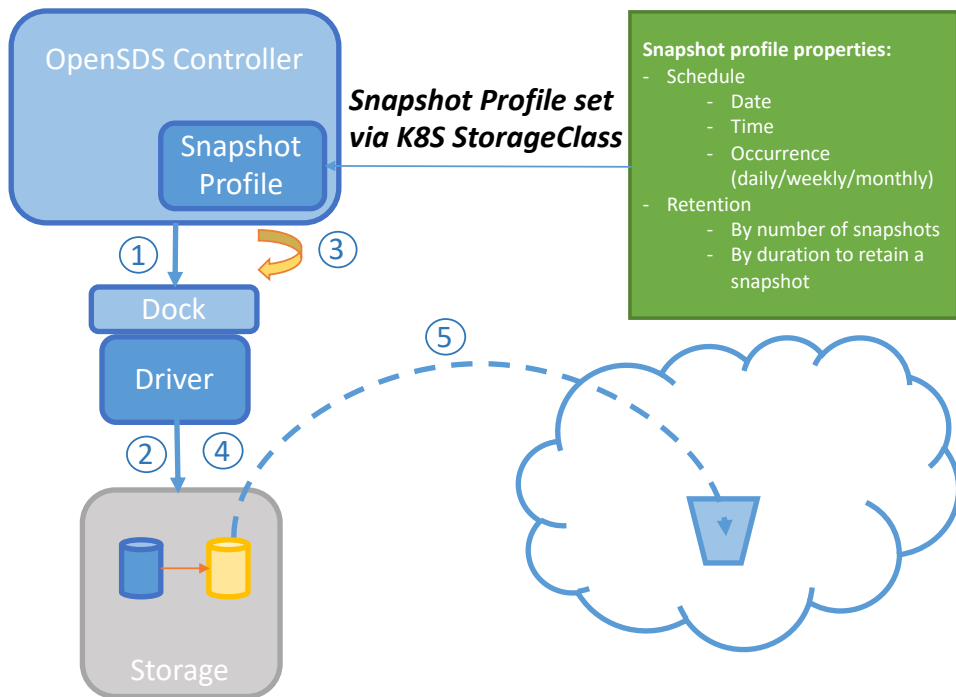
```
docker exec -it <nginx container id> /bin/bash
```

```
root@nginx:/# mount | grep html
/dev/sda on /var/lib/www/html type ext4 (rw,relatime,data=ordered)
```

nginx.yaml

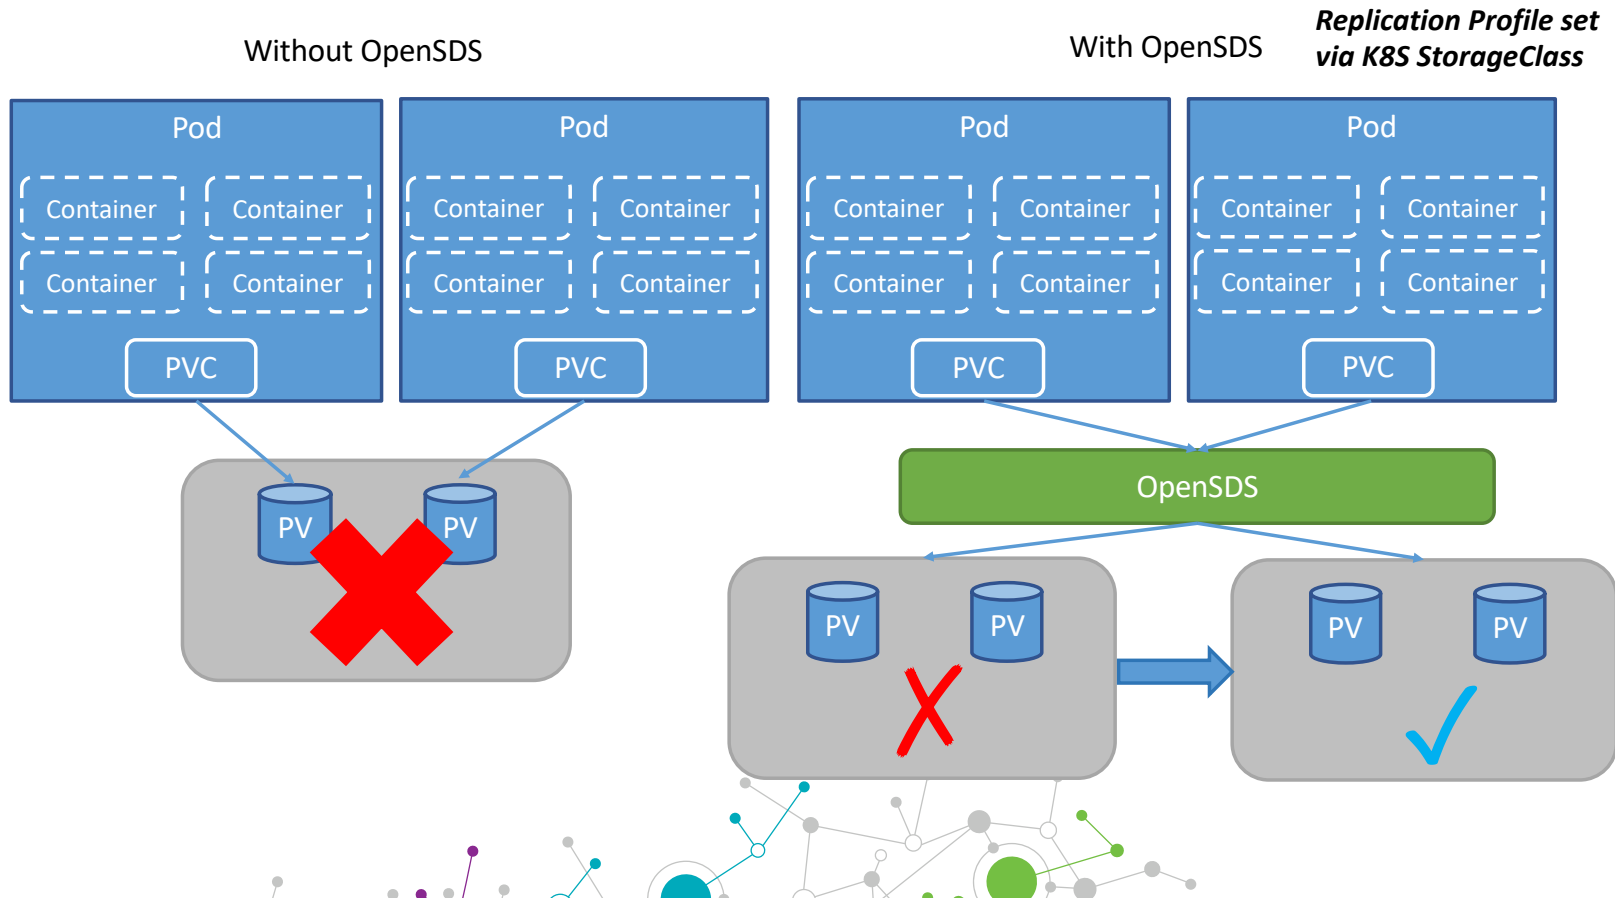
```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - image: nginx
      imagePullPolicy: IfNotPresent
      name: nginx
      ports:
        - containerPort: 80
          protocol: TCP
      volumeMounts:
        - mountPath: /var/lib/www/html
          name: csi-data-opensdsplugin
  volumes:
    - name: csi-data-opensdsplugin
      persistentVolumeClaim:
        claimName: opensds-csi-high-performance-pvc
        readOnly: false
```

# Data Protection for Persistent Volumes




1. Controller asks driver to create a volume.
2. Driver creates a volume on the storage backend.
3. Controller periodically asks driver to create a snapshot based on policies defined in the Snapshot Profile.
4. Driver creates a snapshot on the storage backend.
5. Driver uploads the snapshot to an object store on premise or in the cloud based on the snapshot profile.

# Disaster Recovery for Persistent Volumes






# Array-based Replication



1. Creates source volume
  - Creates entry in db
  - Creates volume on Storage1.
2. Creates target volume
  - Creates entry in db
  - Creates volume on Storage2
3. Creates source replication
  - Creates entry in db
  - Creates replication relationship on Storage1 and Storage2
4. Controller 1 communicates with controller 2 to create target replication
5. Controller 2 creates entry in db

# Host-based Replication



- Creates source volume
  - Creates entry in db
  - Creates volume on Storage1
- Creates target volume
  - Creates entry in db
  - Creates volume on Storage2
- Attach source volume to Host1
  - Update volume entry in db with host info
- Attach target volume to Host2
  - Update volume entry in db with host info
- Controller 1 Creates source replication
  - Creates entry in db
  - Creates replication relationship on Host1 and Host2 (Host1 is primary)
- Controller 1 communicates with controller 2 to create target replication
- Controller 2 creates entry for target replication in db

# Replication Functionalities

## Create Replication:

`osdsctl replication create <primary volume id> <secondary volume id> [flags]`

## Flags:

<code>-d, --description string</code>	the description of created replication
<code>-h, --help</code>	help for create
<code>-n, --name string</code>	the name of created replication
<code>-p, --primary_driver_data string</code>	the primary replication driver data of created replication
<code>-m, --replication_mode string</code>	the replication mode of created replication, value can be sync/async
<code>-t, --replication_period int</code>	the replication period of created replication, the value must be greater than 0
<code>-s, --secondary_driver_data string</code>	the secondary replication driver data of created replication

## Enable Replication:

`osdsctl replication enable <replication id>`

## Disable Replication:

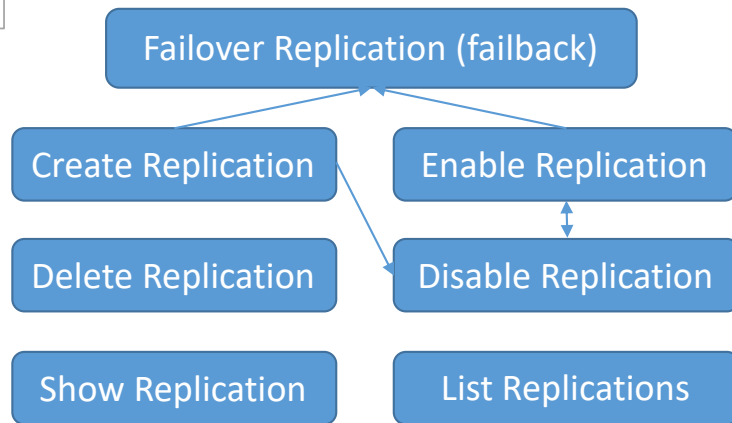
`osdsctl replication disable <replication id>`

## Failover Replication:

`osdsctl replication failover <replication id> [flags]`

## Flags:

<code>-a, --allow_attached_volume</code>	whether allow attached volume when failing over replication
<code>-h, --help</code>	help for failover
<code>-s, --secondary_backend_id string</code>	the secondary backend id of failover replication



# Demo

- Provision storage using OpenSDS CSI plugin based on profile definition
- Create tenant, create profile, create volume and snapshot via UI
- Array-based replication using Dorado
- Note: Host-based replication using DRBD is WIP



# Future Development

- Array-based and Host-based Migration
- Multi-Cloud Storage
- Multi-OpenStack



# Join Us

- Repos: <https://github.com/opensds>
- Slack: [opensds.slack.com](https://opensds.slack.com)
- Mailing list: [opensds-tech-discuss@lists.opensds.io](mailto:opensds-tech-discuss@lists.opensds.io)



# Thank You

@opensds\_io